How to use Pandas for SQL-like actions?

Hessam Mohammad Hosseini

January 3, 2023

Abstract

This document is part 1 of my cheat sheet on **Pandas** which provide overall review on how to use **Pandas** for those familiar with **SQL**.

Contents

In this cheetsheet, I try to make readable and easy to use reference for SQL users aiming to explain how to do similar action in Pandas. My assumption is you know how to have your data as dataframe in Pandas. As soon as you have a dataframe, you can query like a table in SQL. Many possibilities are avaiable IO Tools (CSV, EXCEL,DB connection ...)but most important ones are reading from Excel and CSV.

For example, following line of code will read file located at `file_address`, skip first row, consider `Time` column to be parsed as `datetime` format and consider, as thousands identifer while reading file.

```
df = pd.read_csv(file_name,
                 skiprows=10, # ignore first 10 rows of file
                 parse_dates=['Time'], # Parse Time as datetime column
                 thousands=',') # numbers have , to demonstrate 000 sepration
```

1. read_csv Function has lots of useful options to fasilitate work and avoid doing extra data cleaning tasks after data loading. Here are options for `read_csv`.
   ```
   df = pandas.read_csv(filepath_or_buffer, sep=', ', delimiter=None,
   header='infer', names=None, index_col=None, usecols=None,
   squeeze=False, prefix=None, mangle_dupe_cols=True, dtype=None,
   engine=None, converters=None, true_values=None, false_values=None,
   skipinitialspace=False, skiprows=None, skipfooter=0, nrows=None,
   na_values=None, keep_default_na=True, na_filter=True, verbose=False,
   skip_blank_lines=True, parse_dates=False, infer_datetime_format=False,
   keep_date_col=False, date_parser=None, dayfirst=False, iterator=False,
   chunksize=None, compression='infer', thousands=None, decimal=b'.',
   lineterminator=None, quotechar='"', quoting=0, doublequote=True,
   escapechar=None, comment=None, encoding=None, dialect=None,
   tupleize_cols=None, error_bad_lines=True, warn_bad_lines=True,
   delim_whitespace=False, low_memory=True, memory_map=False,
   float_precision=None)
   ```
2. read_excel Functions has lots of useful options to fasilitate work. Here are options for `read_excel`:
   ```
   df = pandas.read_excel(io, sheet_name=0, header=0, names=None,
   index_col=None, parse_cols=None, usecols=None, squeeze=False, dtype=None,
   engine=None, converters=None, true_values=None, false_values=None,
   skiprows=None, nrows=None, na_values=None, keep_default_na=True,
   verbose=False, parse_dates=False, date_parser=None, thousands=None,
   comment=None, skip_footer=0, skipfooter=0, convert_float=True,
   mangle_dupe_cols=True, **kwds)
   ```

I tried to summerized and add to what was available on Pandas comparison with SQL aiming to simplify understanding. For details on functionality, please check and review Pandas documentation.

1 SELECT

SQL Sample	Pandas Sample
`select * FROM table`	`table`
`select distinct c5 FROM table`	`table['c5'].unique()` or `table[['c5']].drop_duplicates()`
`select c1, c2,c10 FROM table`	`df[[c1, c2,c10]]`
`select c10, c2, c1 FROM table`	`df[[c10, c1, c2]]`
`select c10, c2*12 - c1 + c6 FROM table`	`df['new c'] = df.c2*12 - df.c1 + df.c6df[[c10,'new c']]` or `df.assign(c10 = df.c10, new_c = df.c2*12 - df.c1 + df.c6)`
`SELECT total_bill, tip, smoker, time FROM tips LIMIT 5`	`tips[['total_bill', 'tip', 'smoker','time']].head(5)`

1.1 Update or delete

	SQL Sample	Pandas Sample
Delete	`DELETE FROM tips WHERE tip > 9`	`tips = tips.loc[tips['tip'] <= 9]`
Update	`UPDATE tips SET tip = tip*2 WHERE tip < 2`	`tips.loc[tips['tip'] < 2, 'tip'] *= 2`

2 Conditioning at WHERE

	SQL Sample	Pandas Sample
-	`SQL SELECT * FROM tips WHERE time = 'Dinner' LIMIT 5`	`tips[tips['time'] == 'Dinner'].head(5)` or `is_dinner = tips['time'] == 'Dinner'tips[is_dinner].head(5)`
AND	`SELECT * FROM tips WHERE time = 'Dinner' AND tip > 5.00`	`tips[(tips['time'] == 'Dinner') & (tips['tip'] > 5.00)]`

	SQL Sample	Pandas Sample
OR	SELECT * FROM tips WHERE size >= 5 OR total_bill > 45	tips[(tips['size'] >= 5') or (tips['total_bill'] > 45)]
IS NULL	SELECT * FROM t WHERE col2 IS NULL	t[t['col2'].isna()]
IS NOT NULL	SELECT * FROM t WHERE col IS NOT NULL	t[t['col2'].notna()]

SQL Sample	Pandas Sample
SELECT * FROM tips ORDER BY tip DESC LIMIT 10 OFFSET 5	tips.nlargest(10 + 5, columns='tip').tail(10)
SELECT total_bill, tip, smoker, time FROM tips ORDER BY tip DESC LIMIT 10 OFFSET 5	tips[['total_bill', 'tip', 'smoker','time']] tips.nlargest(10 + 5, columns='tip').tail(10)

3 GROUP BY

SQL Sample	Pandas Sample
SELECT sex, count(*) FROM tips GROUP BY sex select A, sum(C), sum(D) FROM df GROUP BY A	tips.groupby('sex').size() or tips.groupby('sex')['total_bill'].count() df.groupby('A')['B','C'].sum()

SQL Sample	Pandas Sample
SELECT day, AVG(tip), COUNT(*) FROM tips GROUP BY day	tips.groupby('day').agg({'tip': np.mean, 'day': np.size})

SQL Sample	Pandas Sample
SELECT smoker, day, COUNT(*), AVG(tip) FROM tips GROUP BY smoker, day	tips.groupby(['smoker','day']).agg({'tip': [np.size, np.mean]})

SQL Sample	Pandas Sample
SELECT c1, COUNT(*) FROM df where country='IR' GROUP BY c1 having count(*)>1000	df[df.country == 'IR'].groupby('c1').filter(lambda g: len(g) > 1000).groupby('c1').size()
SELECT c1, COUNT(*) FROM df WHERE country='IR' GROUP BY c1 HAVING count(*)>1000 ORDER BY count(*) desc	df[df.country == 'IR'] .groupby('c1'). filter(lambda g: len(g) > 1000) .groupby('c1').size() .sort_values(ascending=False)

4 ORDER BY

SQL Sample	Pandas Sample
SELECT * FROM df ORDER BY A, B	df.sort_values(['A', 'B'])
SELECT * FROM df ORDER BY A desc, C	df.sort_values(['A', 'B'],ascending=[False, True])

5 UNION, JOIN and other set related operations

I will work to provide more comprehensive explanations on this part.

5.1 Union

SQL Sample	Pandas Sample
SELECT c1, c2 FROM df1 UNION ALL SELECT c1, c2 FROM df2	pd.concat([df1, df2])

Difference between `union all` and `union` is that `union` will remove duplicates.

SQL Sample	Pandas Sample
SELECT c1, c2 FROM df1 UNION SELECT c1, c2 FROM df2	pd.concat([df1, df2]) .drop_duplicates()

5.2 Different Join cases

I will add parts to make eplanation on `join` more clear and compehensive. Below image extracted from Enthoughtnamed "Enthought-Python-Pandas-Cheat-Sheets-1-8-v1.0.2" worth more than 100 sentences to explain different types of `join`. You can get whole file via registration on Enthought.

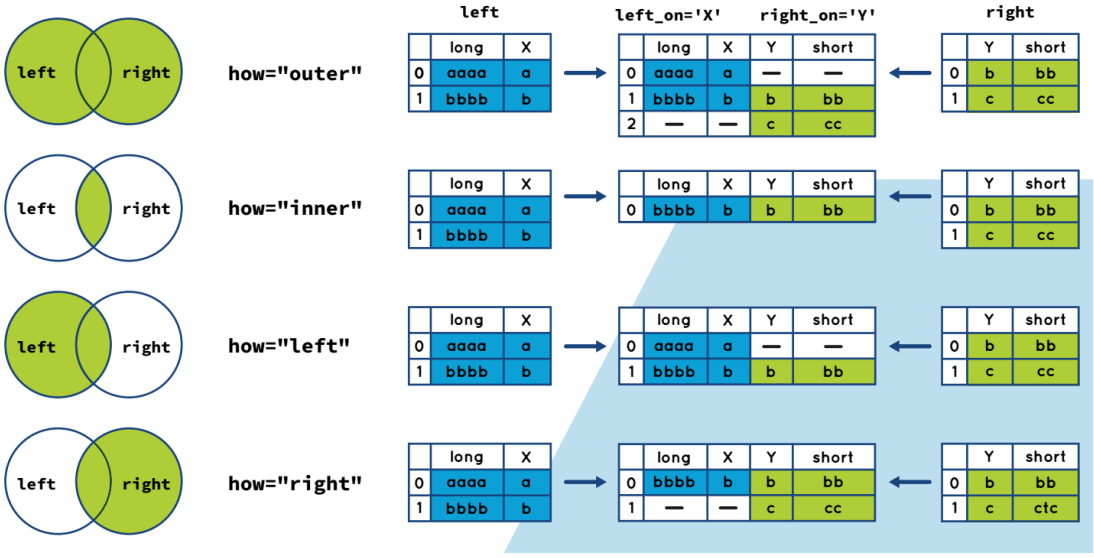

SQL Sample	Pandas Sample
SELECT * FROM df1 INNER JOIN df2 ON df1.key = df2.key	pd.merge(df1, df2, on='key')
SELECT * FROM df1 INNER JOIN df2 ON df1.c7 = df2.c5	pd.merge(df1, df2, left_on='c7',right_on='c5')

6 Time functionality

In order to have possibility to use time related functionalities, we need help `Pandas` understand which columns are to be treated as time. Of course, the columns should be in for that converting them to time format is possible. For details, please check and review Pandas documentation. If you manage to let `Pandas` know properly which column(s) to be time related column(s), they will end up having `datetime64[ns]` format. `.dtypes` on dataframe provides you with columns formats. Pass date related column(s) you need to `parse_dates` to `read_csv` or `read_excel` functions. Check Pandas documentation for more details.

Doing so, you can apply `.dt` on column to have date - time selection like

`dt.dayofweek`

`dt.minute`

`dt.hour`

`dt.second`

`dt.quarter`

`dt.month`

`dt.month_name`

`dt.weekday_name`

`dt.weekday`

`dt.weekofyear`

`dt.year`

6.1 How to get current date time using pandas?

`pd.datetime.now()`

`pd.datetime.now().date()`

`pd.datetime.now().year`

`pd.datetime.now().month`

`pd.datetime.now().day`

`pd.datetime.now().hour`

`pd.datetime.now().minute`

`pd.datetime.now().second`

`pd.datetime.now().microsecond`

Again, check Pandas documentation for more! Here, we assume `sdate` column to have `datetime64[ns]` format.

	SQL Sample	Pandas Sample
sysdate - n	`SELECT * FROM df WHERE sdate> sysdate-5`	`df[df['sdate'].dt.date()> pd.datetime.now().date()-5]`
month	`SELECT * FROM df WHERE sdate in Q1`	`df[(df.sdate.dt.month >= 1) & (df.sdate.dt.month <= 3)]`

	SQL Sample	Pandas Sample
between	SELECT * FROM t WHERE to_char(sdate,'yyyy') between 1998 AND 2018	t[(t.sdate.dt.year >= 1998) & (t.sdate.dt.year <= 2018)]
	SELECT * FROM t WHERE to_char(sdate ,'day')= 'Friday'	df[df.sdate.dt.day_name() == 'Friday']

7 String related functionality like `like`, `Substr`

For columns with `string` content, we could access string related funftionality by applying `.str` on column. Here are few samples: `str.contains` - contains options:

```
Series.str.contains(pat, case=True, flags=0, na=nan, regex=True)
```

Here are list of main ‘string‘ functions.

```
str.upper
str.lower
str.extract
str.extractall
str.find
str.findall
str.len
str.replace
str.slice
str.split
str.strip
```

Check Pandas documentation for more!

	SQL Sample	Pandas Sample
regex	SELECT upper(trim(to_char(LAC,'xxxxx')) ||'-' || trim(to_char(CI,'xxxxx'))) AS "LAC-CI(HEX)" FROM t	t = t['LAC','CI']\ .apply(lambda x: x\ .astype(str)\ .map(lambda x: int(x, base=16)))t .assig(LAC-CI(HEX) = t['LAC']+'-'+t['CI']
substr	SELECT * FROM tips WHERE substr(time,1,2) = 'Di'	tips[tips['time']\ .str[:2] == 'Di']

	SQL Sample	Pandas Sample
regex	SELECT upper(trim(to_char(LAC,'xxxxx')) ||'-' || trim(to_char(CI,'xxxxx'))) AS "LAC-CI(HEX)" FROM t	t = t['LAC','CI']\ .apply(lambda x: x\ .astype(str)\ .map(lambda x: int(x, base=16)))t .assig(LAC-CI(HEX) = t['LAC']+'-'+t['CI']
like	SELECT * FROM df WHERE Country like '%IR%'	df[df['Country']\ .str.contains('IR') == True]
like	SELECT * FROM df WHERE Country like 'IR%'	df[df['Country']\ .str.startswith('IR') == True]
like	SELECT * FROM df WHERE Country like '%AN'	df[df['Country']\ .str.endswith('AN') == True]
in	SELECT * FROM df WHERE City in ('TEHRAN', 'BERLIN','STOKHOLM')	df[df['City']\ .isin(['TEHRAN', 'BERLIN','STOKHOLM'])
regex	SELECT last_name FROM contacts WHERE REGEXP_LIKE (last_name, '^+A(*)')	contacts[contacts['last_name']\ .str.contains('^+A(*)')]
regex	SELECT c1 FROM t WHERE REGEXP_LIKE(c1,'([A-Z][\d]{4})')	t[t['c1']\ .str.contains(([A-Z][\d]{4}))]

|| provide concatenation functonality in PL/SQL. In **Python**, + on string values resulted in contanated text.

SQL Sample	Pandas Sample
SELECT * FROM (SELECT t.*, ROW_NUMBER() OVER(PARTITION BY day ORDER BY total_bill DESC) AS r FROM tips t) WHERE r < ORDER BY day, r;	(tips.assign(r=tips\ .sort_values(['total_bill'], ascending=False)\ .groupby(['day'])\ .cumcount()+1)\ .query('r < 3') .sort_values(['day', 'r']))

Oracle's ROW_NUMBER() analytic function

8 To check for missing values

- `df.notnull()` Use to Drop Missing Values
- `df.dropna()` Filling Missing Values — Direct Replace
- `df.fillna()`

Besides Pandas comparison with SQL, I also get ideas from following references:

1. pandas-cheatsheet-for-sql-people-part-1
2. how-to-rewrite-your-sql-queries-in-pandas-and-more
3. thinking-like-sql-in-pandas
4. did-you-know-pandas-can-do-so-much
5. 10-python-pandas-tricks-that-make-your-work-more-efficient